

Introduction to Computer Vision
Problem Set 1

2. d) I used the default accumulator bin size of 1, because the image changed very sharply so I needed 1 pixel of accuracy.

I also used the default threshold of $0.5 * \max$ because there were enough points on the lines to outweigh any faulty lines.

I also used the default neighborhood size of $\text{size}(H)/50$ because the actual lines in the image were pretty far apart.